

**6. To design and deploy a multi-container application using Docker Compose by creating an application service and a dependent database/Redis service, configuring service dependencies, networks, and volumes, and verifying inter-container communication.**

#### **Prerequisites (Ubuntu)**

Make sure Docker and Docker Compose are installed.

```
bash  
  
docker --version  
docker compose version
```

If not installed (usually not needed in exam, but for completeness):

```
bash  
  
sudo apt update  
sudo apt install docker.io docker-compose-plugin -y  
sudo systemctl start docker  
sudo systemctl enable docker
```

#### **Step 1: Create a project directory**

```
bash  
  
mkdir multi_container_app  
cd multi_container_app
```

#### **Step 2: Create application code (simple Python app)**

This app connects to Redis and stores a key.

Create a file `app.py`:

```
bash  
  
nano app.py
```

Paste this code:

```
python  
  
import redis  
import time  
  
time.sleep(5) # wait for redis to start  
  
r = redis.Redis(host='redis', port=6379, decode_responses=True)  
r.set('message', 'Hello from App Container')  
  
print("Stored message in Redis:", r.get('message'))
```

#### **Step 3: Create requirements file**

```
bash  
  
nano requirements.txt  
  
text  
  
redis
```



## Step 4: Create Dockerfile for application service

```
bash

nano Dockerfile

dockerfile

FROM python:3.10-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install -r requirements.txt

COPY app.py .

CMD ["python", "app.py"]
```

## Step 5: Create Docker Compose file

This is the **core part** of the lab answer.

```
bash

nano docker-compose.yml

yaml

version: "3.8"

services:
  app:
    build: .
    depends_on:
      - redis
    networks:
      - app_network

  redis:
    image: redis:latest
    volumes:
      - redis_data:/data
    networks:
      - app_network

networks:
  app_network:

volumes:
  redis_data:
```

### **Step 6: Run the multi-container application**

```
bash  
  
docker compose up --build
```

### **Step 7: Verify inter-container communication**

Expected output:

```
text  
  
Stored message in Redis: Hello from App Container
```

### **Step 8: Verify running containers (optional but good)**

```
bash  
  
docker ps
```

### **Step 9: Verify volume creation**

```
bash  
  
docker volume ls  
  
You should see:  
  
text  
  
redis_data
```

### **Step 10: Stop containers**

```
bash  
  
docker compose down
```

7. To implement data ingestion, cleaning, and versioning using Data Version Control (DVC) by tracking raw and processed datasets, creating a reproducible pipeline with dvc.yaml, and validating reproducibility via dvc repro.

## Step 1: Install prerequisites (Ubuntu)

```
bash

sudo apt update
sudo apt install git python3 python3-pip -y
pip3 install dvc
```

Verify:

```
bash

dvc --version
```

## Step 2: Create project directory

```
bash

mkdir dvc_lab
cd dvc_lab
```

## Step 3: Initialize Git and DVC

```
bash

git init
dvc init
```

## Step 4: Create data directories

```
bash

mkdir data
mkdir data/raw data/processed
```

## Step 5: Create raw dataset (Data Ingestion)

Create a simple CSV file.

```
bash

nano data/raw/data.csv

c

id,name,age
1,Alice,23
2,Bob,
3,Charlie,25
```

## Step 6: Track raw data using DVC

```
bash

dvc add data/raw/data.csv
git add data/raw/data.csv.dvc .gitignore
git commit -m "Add raw dataset with DVC tracking"
```

### Step 7: Create data cleaning script

```
bash
nano clean_data.py

python
import pandas as pd

df = pd.read_csv("data/raw/data.csv")

# Remove rows with missing values
df = df.dropna()

df.to_csv("data/processed/cleaned_data.csv", index=False)
```

Install dependency:

```
bash
pip3 install pandas
```

### Step 8: Create DVC pipeline (dvc.yaml)

This is the core requirement of the question.

```
bash
dvc run -n clean_data \
-d clean_data.py \
-d data/raw/data.csv \
-o data/processed/cleaned_data.csv \
python3 clean_data.py
```



### Step 9: Track pipeline files

```
bash
git add dvc.yaml dvc.lock clean_data.py
git commit -m "Add data cleaning pipeline"
```

### Step 10: Track processed data

```
bash
dvc add data/processed/cleaned_data.csv
git add data/processed/cleaned_data.csv.dvc
git commit -m "Track processed dataset"
```

### Step 11: Verify reproducibility using `dvc repro`

```
bash
dvc repro
```

Expected output:

```
text
Stage 'clean_data' didn't change, skipping
```

## Step 12: Validate versioning (optional but good)

Modify raw data:

```
bash

nano data/raw/data.csv
```

Add a new row:

```
c

4,David,28
```

Run:

```
bash

dvc repro
```

8. To perform experiment tracking using MLflow by training a machine learning model, logging metrics, parameters, and artifacts, and comparing multiple runs to identify the best-performing model.

### Step 1: Install required packages (Ubuntu)

```
bash

sudo apt update
sudo apt install python3 python3-pip -y
pip3 install mlflow scikit-learn pandas
```

Verify:

```
bash

mlflow --version
```

### Step 2: Create project directory

```
bash

mkdir mlflow_lab
cd mlflow_lab
```

### Step 3: Create training script

Create `train.py`:

```
bash

nano train.py
```

```
python
```

```
Copy code
```

```
import mlflow
import mlflow.sklearn
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load dataset
X, y = load_iris(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Different hyperparameter values
C_value = 1.0

with mlflow.start_run():
    model = LogisticRegression(C=C_value, max_iter=200)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)

    # Log parameters
    mlflow.log_param("C", C_value)
    mlflow.log_param("model", "LogisticRegression")

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)

    # Log model as artifact
    mlflow.sklearn.log_model(model, "model")

    print("Accuracy:", accuracy)
```

#### Step 4: Run first experiment

```
bash
python3 train.py
```

#### Step 5: Run second experiment with different parameter

Edit `train.py`:

```
bash
nano train.py
```

Change:

```
python
C_value = 0.1
```

Run again:

```
bash
python3 train.py
```

Now you have **multiple runs** logged in MLflow.

#### Step 6: Start MLflow Tracking UI

```
bash
mlflow ui
```

Open browser:

```
text
http://127.0.0.1:5000
```

9. To optimize and standardize model inference using ONNX by exporting a trained ML model toONNX format, running inference with ONNX Runtime, and benchmarking performance against the native scikit-learn model.

#### Step 1: Install required packages (Ubuntu)

```
bash
sudo apt update
sudo apt install python3 python3-pip -y
pip3 install scikit-learn onnx onnxruntime skl2onnx numpy
```

Copy code

Verify:

```
bash
python3 -c "import onnx, onnxruntime; print('ONNX setup OK')"
```

Copy code

#### Step 2: Create project directory

```
bash
mkdir onnx_lab
cd onnx_lab
```

Copy code

#### Step 3: Create Python script

```
bash
nano onnx_inference.py
```

Copy code

```
import time
import numpy as np
import onnxruntime as ort

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from skl2onnx import convert_sklearn
from skl2onnx.common.data_types import FloatTensorType

# Load dataset
X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Train scikit-learn model
model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)

# -----
# Native scikit-learn inference
# -----
start = time.time()
sk_preds = model.predict(X_test)
sk_time = time.time() - start

print("Scikit-learn inference time:", sk_time)

# -----
# Export model to ONNX
# -----
initial_type = [("input", FloatTensorType([None, X.shape[1]]))]
onnx_model = convert_sklearn(model, initial_types=initial_type)

with open("model.onnx", "wb") as f:
    f.write(onnx_model.SerializeToString())

print("Model exported to ONNX format")

# -----
# ONNX Runtime inference
# -----
session = ort.InferenceSession("model.onnx", providers=["CPUExecutionProvider"])
input_name = session.get_inputs()[0].name

start = time.time()
onnx_preds = session.run(
    None, {input_name: X_test.astype(np.float32)}
)
onnx_time = time.time() - start

print("ONNX Runtime inference time:", onnx_time)
```

#### Step 4: Run the program

```
bash
python3 onnx_inference.py
```

#### Step 5: Expected output (example)

```
text
Scikit-learn inference time: 0.0012
Model exported to ONNX format
ONNX Runtime inference time: 0.0004

(Exact timings may vary, but ONNX is usually faster.)
```

#### Step 6: Verify ONNX model file

```
bash
ls
You should see:
text
model.onnx  onnx_inference.py
```

10. To serve a machine learning model using FastAPI by developing a /predict REST endpoint with input validation, writing test cases, containerizing the service, and verifying predictions through Postman or curl.

#### Step 1: Install required packages (Ubuntu)

```
bash
sudo apt update
sudo apt install python3 python3-pip docker.io -y
pip3 install fastapi uvicorn scikit-learn pydantic pytest
```

Start Docker:

```
bash
sudo systemctl start docker
sudo systemctl enable docker
```

#### Step 2: Create project directory

```
bash
mkdir fastapi_ml
cd fastapi_ml
```

#### Step 3: Create ML model and API

Create `main.py`:

```
bash
nano main.py
```

```
python

from fastapi import FastAPI
from pydantic import BaseModel
from sklearn.datasets import load_iris
from sklearn.linear_model import LogisticRegression

# Load and train model
X, y = load_iris(return_X_y=True)
model = LogisticRegression(max_iter=200)
model.fit(X, y)

app = FastAPI()

# Input validation
class IrisInput(BaseModel):
    sepal_length: float
    sepal_width: float
    petal_length: float
    petal_width: float

@app.post("/predict")
def predict(data: IrisInput):
    features = [
        data.sepal_length,
        data.sepal_width,
        data.petal_length,
        data.petal_width
    ]
    prediction = model.predict(features)
    return {"prediction": int(prediction[0])}
```

#### Step 4: Run FastAPI locally

bash

```
uvicorn main:app --reload
```

Test in browser:

text

```
http://127.0.0.1:8000/docs
```

## Step 6: Write test case

Create `test_api.py`:

```
bash

nano test_api.py

python

from fastapi.testclient import TestClient
from main import app

client = TestClient(app)

def test_predict():
    response = client.post("/predict", json={
        "sepal_length": 5.1,
        "sepal_width": 3.5,
        "petal_length": 1.4,
        "petal_width": 0.2
    })
    assert response.status_code == 200
    assert "prediction" in response.json()
```

Run test:

```
bash

pytest
```

## Step 7: Create Dockerfile (containerization)

```
bash

nano Dockerfile

dockerfile

FROM python:3.10-slim

WORKDIR /app

COPY main.py .
RUN pip install fastapi uvicorn scikit-learn pydantic

CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

## Step 8: Build and run Docker container

```
bash

sudo docker build -t fastapi-ml .
sudo docker run -p 8000:8000 fastapi-ml
```

## What you should do instead (container case)

### ✓ Alternative: Swagger UI inside the container

After running the container:

```
bash
sudo docker run -p 8000:8000 fastapi-ml
```

Open in browser:

```
bash
http://localhost:8000/docs
```

Then:

1. Select POST /predict
2. Click Try it out
3. Enter input JSON
4. Click Execute
5. See prediction output

This verifies the containerized API completely.

1. Create a simple User Registration Form for an event with fields (Name, Email, Phone). Initialize a Git repository, commit the project files, and push them to GitHub using appropriate Git commands. Create a new branch named update-form, modify the registration form by adding a new field (Department), merge the branch into main, and push the changes to GitHub. Show the use of: git branch, git checkout, git merge, git push.

```
mkdir student-portal
cd student-portal
nano index.html
```

```
<!doctype html>
<html lang="en">
<head>
<meta charset="utf-8"/>
<meta name="viewport" content="width=device-width,initial-scale=1"/>
<title>Event Registration</title>
</head>
<body>
<h1>Event Registration</h1>
<form id="regForm" action="#" method="post">
<label>Name: <input type="text" name="name" required></label><br>
<label>Email: <input type="email" name="email" required></label><br>
<label>Phone: <input type="tel" name="phone" required></label><br>
<button type="submit">Register</button>
</form>
</body>
```

```
</html>
```

```
git init  
git config --global user.name "Your Name"  
git config --global user.email "you@example.com"
```

```
git add index.html  
git commit -m "Initial commit: add registration form"
```

#### **Step 8: Create GitHub repository (MANUAL STEP)**

1. Go to [github.com](https://github.com)
2. Create new repository named:

```
student-portal
```

3. Do NOT add README or license (important)

#### **Step 9: Add remote repository**

(Use HTTPS — easier in exam labs)

```
bash  
git remote add origin https://github.com/YOUR_USERNAME/student-portal.git
```

Verify:

```
bash  
git remote -v
```

```
git branch -M main  
git push -u origin main
```

#### **Step 11: Verification commands (SHOW in exam)**

```
bash  
git log --oneline  
bash  
git branch  
bash  
git status
```

#### **PART B — Create branch, modify form, merge, push**

##### **Step 12: Create and switch to new branch**

```
bash  
git checkout -b update-form  
Verify:  
bash  
git branch
```

**Step 13: Modify index.html (Add Department field)**

```
bash  
  
nano index.html
```

Add this line AFTER Phone:

```
html  
  
<label>Department: <input type="text" name="department" required></label>
```

Save and exit.

**Step 14: Stage and commit changes**

```
bash  
  
git add index.html  
git commit -m "Add Department field to registration form"
```

**Step 15: Switch back to main branch**

```
bash  
  
git checkout main
```

```
git merge update-form  
git push origin main
```

```
git branch -d update-form  
git show HEAD:index.html
```

## 2. Jenkins

**A. Repo (on your machine)**

If you don't already have the repo locally, clone it:

```
bash  
  
git clone https://github.com/YOUR_USERNAME/student-portal.git  
cd student-portal
```

 Copy code

**B. Create the Jenkins Freestyle job (UI steps — exact)**

1. Open [http://<JENKINS\\_HOST>/](http://<JENKINS_HOST>/) → New Item.
2. Enter name: `student-portal-site` → select Freestyle project → OK.
3. Source Code Management → choose Git → Repository URL:  
`https://github.com/YOUR_USERNAME/student-portal.git`
  - If private repo: add credentials (username/password or SSH key) using Add.
4. Build Triggers → check GitHub hook trigger for GITScm polling (or "Build when a change is pushed to GitHub").
5. Build → Add build step → Execute shell. Put these exact lines:

```
bash  
  
echo "Building student-portal"  
ls -la  
echo "Contents of index.html:"  
sed -n '1,80p' index.html
```

 Copy code

6. Save the job.

### C. Configure GitHub webhook (UI steps — exact)

1. On GitHub repo → **Settings** → **Webhooks** → **Add webhook**.

2. Payload URL:

```
http://<JENKINS_HOST>/github-webhook/
```

3. Content type: `application/json`

4. Which events: **Just the push event** → **Add webhook**.

(If Jenkins is not publicly reachable, use `ngrok` or use GitHub Enterprise inside lab — your instructor will usually provide connectivity.)

### D. Make a tiny change and push (exact commands)

Run these inside the cloned repo:

```
bash

# create small branch, change index.html heading, and push
git checkout -b tiny-update
sed -i 's/Event Registration/Event Registration - Updated/' index.html
git add index.html
git commit -m "Tiny update: change heading for CI test"
git push -u origin tiny-update

# merge to main (or create PR and merge); here we merge locally and push
git checkout main
git merge tiny-update
git push origin main
```

 Copy code

**Result:** GitHub receives `push` event → webhook posts to Jenkins → Jenkins starts `student-portal-site` job.



### E. Verify Jenkins build triggered (commands you can run)

1. (Quick) check job console via CLI — if Jenkins has no auth or you have a token:

```
bash

# Get console text of last build
curl "http://<JENKINS_HOST>/job/student-portal-site/lastBuild/consoleText"
```

 Copy code

You should see:

```
php-template

Building student-portal
total 8
-rw-r--r-- ... index.html
Contents of index.html:
<!doctype html>
...
<h1>Event Registration - Updated</h1>
...
Finished: SUCCESS
```

 Copy code

3. Create a simple web application (HTML or Python Flask). Write a Dockerfile, build a Docker image, and run the container so the application is accessible on a browser using port mapping. Using Docker commands, perform the following operations: List images and containers Stop a running container Remove a container and image Demonstrate commands like: `docker ps`, `docker stop`, `docker rm`, `docker rmi`.

```
bash
mkdir student-portal-docker
cd student-portal-docker

Step 2: Create app.py
bash
nano app.py

Paste EXACTLY THIS:
python

from flask import Flask, render_template_string
app = Flask(__name__)

HTML = """
<!doctype html>
<html><body>
<h1>Student Portal</h1>
<p>Registration form:</p>
<form>
Name: <input name="name"><br>
Email: <input name="email"><br>
Phone: <input name="phone"><br>
Department: <input name="department"><br>
</form>
</body></html>
"""

@app.route("/")
def home():
    return render_template_string(HTML)

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=5000)
```



### Step 3: Create requirements.txt

```
bash
nano requirements.txt

text
Flask==2.3.2
```

### Step 4: Create Dockerfile

```
bash
nano Dockerfile

Paste EXACTLY THIS:
dockerfile

FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY app.py .
EXPOSE 5000
CMD ["python", "app.py"]
```

## **Build docker image**

```
docker build -t student-portal:v1 .
docker run -d --name student-portal -p 5000:5000 student-portal:v1
```

Open <http://localhost:5000>

## **Listing running containers**

```
docker ps
```

## **List all containers**

```
docker ps -a
```

## **List docker images**

```
docker images
```

## **Stop running container**

```
docker stop student-portal
```

## **Remove container**

```
docker rm student-portal
```

## **Remove docker image**

```
docker rmi student-portal:v1
```

4. Deploy the previously created Docker image on Kubernetes using a Deployment YAML file. Verify that the pod is running using kubectl commands. Expose the application using a NodePort Service, access it using the node's IP and port, and scale the deployment to 3 replicas using kubectl scale.

The screenshot shows a terminal window with the following content:

```
Check:
bash
kubectl version --client

2 Cluster is running (VERY IMPORTANT)
For minikube:
bash
minikube start

Verify:
bash
kubectl get nodes

You should see 1 node in Ready state.

PART A — Make Docker Image Available to Kubernetes
Step 1: Ensure Docker image exists
(This image was created in previous Docker question)
bash
docker images

You should see:

student-portal    v1
```

### **Step 2: Load image into minikube**

```
minikube image load student-portal:v1
```

```
nano deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: student-portal-deploy
spec:
  replicas: 1
  selector:
    matchLabels:
      app: student-portal
  template:
    metadata:
      labels:
        app: student-portal
    spec:
      containers:
        - name: student-portal
          image: student-portal:v1
          ports:
            - containerPort: 5000
```

#### Step 4: Apply deployment

```
bash
kubectl apply -f deployment.yaml
```

##### Expected:

```
bash
deployment.apps/student-portal-deploy created
```

#### Step 5: Verify deployment

```
bash
kubectl get deployments
```

##### Expected:

```
student-portal-deploy  1/1   1   1   20s
```

#### Step 6: Verify pod is running

```
bash
kubectl get pods -l app=student-portal
```

##### Expected:

```
sql
student-portal-deploy-xxxxxx  1/1   Running
```

**Step 7: View pod logs (proof of running app)**

```
bash  
  
kubectl logs <pod-name>  
  
You should see Flask startup logs.
```

**PART C — Expose Application Using NodePort****Step 8: Create NodePort Service YAML**

```
bash  
  
nano service-nodeport.yaml
```

Paste EXACTLY THIS:

```
yaml  
  
apiVersion: v1  
kind: Service  
metadata:  
  name: student-portal-nodeport  
spec:  
  type: NodePort  
  selector:  
    app: student-portal  
  ports:  
  - protocol: TCP  
    port: 5000  
    targetPort: 5000  
    nodePort: 30080
```

Save and exit.

**Step 9: Apply service**

```
bash  
  
kubectl apply -f service-nodeport.yaml
```

**Step 10: Verify service**

```
bash  
  
kubectl get svc student-portal-nodeport  
  
Expected:  
  
student-portal-nodeport  NodePort  10.96.x.x  <none>  5000:30080/TCP
```

**PART D — Access Application****Step 11: Get Node IP**

If using minikube

```
bash  
  
minikube ip
```

Example:

```
192.168.49.2
```

**Step 12: Access application**

Open browser OR use curl: `curl http://<NODE_IP>:30080`

## Scale deployment

```
kubectl scale deployment student-portal-deploy --replicas=3
```

## Verify scaled pods

```
kubectl get pods -l app=student-portal
```

5. Set up a simple ML project environment by creating a requirements.txt file, installing packages, and verifying environment setup. Document the steps in a Jupyter Notebook and commit it to the Git repository.

### STEP 1 — Create project folder & requirements file

```
bash

mkdir ml_env
cd ml_env

cat > requirements-ml.txt << 'EOF'
numpy==1.26.2
pandas==2.2.2
scikit-learn==1.3.2
jupyterlab==4.1.0
EOF
```

[Copy code](#)

### STEP 2 — Create & activate virtual environment

```
bash

python3 -m venv venv
source venv/bin/activate
```

[Copy code](#)

### STEP 3 — Install packages

```
bash

pip install -r requirements-ml.txt
```

[Copy code](#)

### STEP 4 — Verify installation

```
bash

pip freeze | grep -E "numpy|pandas|scikit-learn|jupyter"
```

[Copy code](#)

## STEP 5 — Open Jupyter Notebook

```
bash  
jupyter lab
```

 Copy code

This will open a browser window.

Inside Jupyter Lab UI:

1. Click File → New → Notebook
2. Select Python 3 (ipykernel)
3. Create the two cells exactly as given:
  - Markdown cell
  - Code cell
4. Save as ml\_setup.ipynb

Then run the code cell → output should match requirements.

---

## STEP 6 — Initialize Git

(inside the same ml\_env folder)

```
bash  
git init  
git add requirements-ml.txt ml_setup.ipynb  
git commit -m "Add ML environment requirements and verification notebook"
```

 Copy code

---

## STEP 7 — Connect to GitHub & push

(Replace <your\_repo\_url> with your actual GitHub repo link)

```
bash  
git remote add origin <your_repo_url>  
git branch -M main  
git push -u origin main
```

 Copy code